

F1

F2

F3

Sales Intelligence Assistant

Pre-conversation intelligence for engaging
Director / VP / Partner-level consulting professionals

~60s

per analysis

\$0.08

per run

21+

cited sources

5 days

recency floor

Sales Intelligence Assistant

Pre-conversation intelligence for engaging Director / VP / Partner-level consulting professionals.

Provide at least one input

LinkedIn profile text (copy About + Experience sections)

Paste the public LinkedIn content here...

Company website URL



https://example.com

Freeform notes (optional)

Anything else you know about the person or context...

>  File upload (coming soon)

Analyse

The problem

Senior consulting professionals can spot generic outreach in two seconds

THE MANUAL WORKFLOW

- Open LinkedIn, read their About + Experience
- Google their name + company for recent news
- Check if they've published or spoken publicly
- Look at the firm's current strategic positioning
- Hypothesise what they care about
- Draft a personalised opener

~30 minutes per candidate. Done by ~10% of BDs.

THE CONSEQUENCE

Most outreach to senior consulting professionals gets ignored

Not because the message is bad, but because nothing in it suggests the sender did genuine homework. The bar for earning a reply is structural: visible specificity, visible recency, visible respect.

Design tree

Eight branches mapped before a single line of code

A

Target user

B2B sales + executive search BD
engaging Director / VP / Partner level

C

Output structure

Nine-section JSON
JTBD + Challenger + Push/Pull framing

E

Data enrichment

google_search grounding (recency)
+ BeautifulSoup About/Careers (positioning)

G

Prompt engineering

Schema-bound output, few-shot anchor
Banned phrase list (no 'I hope this finds you well')

B

Input design

Manual paste + scraper for company URL
Three fields, optional, at least one required

D

AI model & reasoning

Gemini Flash (Stage 1) + Pro (Stage 2)
Extended thinking on the analyst stage

F

Architecture

Streamlit single-page app
No auth, no history, single user

H

Workflow design

Staged loading messages
Iterative refine with focus options

Three trade-offs worth defending

Where the design could have gone differently and why it didn't

1

Manual paste over LinkedIn scraping

THE ALTERNATIVE

Alternative: scrape LinkedIn or buy Apollo / PDL enrichment

WHY WE DIDN'T

ToS-compliant, zero data acquisition cost, fully auditable. Production version swaps in Apollo / PDL behind the same UserInputs interface — the prototype should not depend on legally grey infrastructure.

2

Two-stage pipeline over a mega-prompt

THE ALTERNATIVE

Alternative: single Pro prompt with grounding + analysis combined

WHY WE DIDN'T

Reason isn't latency, it's hallucination control. The analyst has no tools — it literally cannot Google. So it cannot invent a citation. Tool isolation is the anti-hallucination spine.

3

Three frameworks layered, not one

THE ALTERNATIVE

Alternative: pick JTBD or Challenger or Push/Pull and apply it cleanly

WHY WE DIDN'T

Headhunting context specifically requires push/pull. Challenger differentiates from generic JTBD. The three layered in order produce structurally different output from any single-framework approach.

Two-stage architecture

Separation of evidence-gathering from reasoning is the spine of everything else



The analyst has no tools — so it can only reason from facts that exist in the sheet. The system literally cannot fabricate a source.

Three-framework analysis stack

Applied in this order, every time

01

Jobs-To-Be-Done

Clayton Christensen

USED FOR

Commercial priorities

SAMPLE OUTPUT

"They are hired to lead complex M&A integrations, so they care about access to high-visibility deal flow."

02

Challenger Sale

Dixon & Adamson

USED FOR

Conversation angles

SAMPLE OUTPUT

"As AI automates analytics, value shifts from the 'how' of integration to the 'who' of counsel."

03

Push / Pull

Executive search tradition

USED FOR

Motivation hypotheses

SAMPLE OUTPUT

*Push: compensation model in transition.
Pull: practice committed to advisory work.*

Anti-hallucination spine

Four layers, in increasing strength

1

Schema enforcement

Stage 2 returns structured output bound to a Pydantic schema. No prose escape hatch.

2

Evidence-ref requirement

Every priority, pain point, and motivation must cite at least one fact_id. Empty list is forbidden.

3

Tool isolation

Stage 2 has no tools. It literally cannot Google. It cannot invent a source it does not see.

4

Data gap surfacing

When facts are insufficient, the analyst MUST populate `data\_gaps` with concrete suggestions instead of producing low-quality content.

F1 — LinkedIn Experience section

Partner at McKinsey & Company.

F2 — LinkedIn Experience section

Joined April 2017, Partner for 9 years 2 months.

F4 — LinkedIn Experience section

Global Co-Lead for Mergers & Acquisitions for Consumer and Retail.

F5 — LinkedIn Experience section

Global Co-Lead for Mergers & Acquisitions for all industry Domains.

F6 — LinkedIn Experience section

lead partner for over 60 large transactions.

F9 — LinkedIn Experience section (A.T. Kearney)

Primary service focus was Mergers & Acquisitions.

F13 — Freeform notes (About Rodrigo)

fluent in English and Spanish and conversant in Portuguese.

F14 — Freeform notes (About Rodrigo)

grew up in Argentina...worked in many countries across the Americas.

F15 — Freeform notes (About Rodrigo)

supported more than 50 mergers and acquisitions...since 2010.

F16 — Google Search / Media Reports (May 2026)

revise its profit-sharing structure...reduce the cash component.

F17 — Google Search / Media Reports (May 2026)

aims to enhance the firm's flexibility and facilitate future investments.

F18 — Google Search / Media Reports (May 2026)

prioritize performance-based pay, largely driven by the impact of AI.

F19 — Google Search / Media Reports (May 2026)

Prompt engineering at a glance

Two prompts, one anchor, and three regeneration directives

STAGE 1	STAGE 2	ANCHOR	REGEN
Research Agent WORD COUNT ~440 WHAT IT DOES Collect facts. No analysis. Every fact must have a source.	Sales Analyst WORD COUNT ~620 WHAT IT DOES Apply JTBD → Challenger → Push/Pull. Cite or omit. No invented facts.	Few-shot example WORD COUNT ~660 WHAT IT DOES One full fictional case (Alex Tanaka) baked into the analyst prompt to anchor tone and depth.	Focus directives WORD COUNT ~85 WHAT IT DOES Three short directives appended on iterative refinement: career motivations / commercial impact / objection handling.

The next six slides walk through these prompts in full, segmented for readability. Source file: prompts.py.

Stage 1 — Research Agent · Part 1 of 2

Role, inputs, rules, and search directives

You are a meticulous research analyst supporting a business development team at an executive search / consulting firm. Their job is to engage senior consulting professionals (Director, VP, Partner level) about career moves or strategic collaborations.

YOUR ROLE

You produce a STRUCTURED FACT SHEET. You do NOT analyse, opine, predict, or recommend. You collect, verify, and label facts with sources.

INPUTS YOU WILL RECEIVE

1. LinkedIn profile text (pasted by the user) – treat as ground truth for stated role/tenure but NOT for inferred intent.
2. Company website content (already scraped from About + Careers pages) – use verbatim, do not infer beyond what is written.
3. Freeform notes from the user – context only.
4. Live Google Search results (you have the google_search tool) – use to find: company news, financial events, product launches in the last 6 months; the person's public talks, articles, podcast appearances, press mentions; any noteworthy controversies or industry awards.

RULES

- Every Fact must have a source_type and source_detail. NO unsourced facts.
- If something is plausible but not verified, DO NOT include it. Leave it out.
- Prefer recent sources (within 6 months) for company news. Older OK for biographical facts.
- Statements must be terse and verifiable: 'Joined Bain in 2018 as Senior Consultant' – not 'Has significant Bain experience'.
- For the 'coverage_assessment' field, be brutally honest. If you only have the LinkedIn headline, say so. The downstream analyst depends on knowing what you DIDN'T find.

WHAT TO SEARCH FOR EXPLICITLY

Run searches for:

```
'[Company name] news [current year]'  
'[Company name] hiring OR layoffs OR expansion'  
'[Person name] [Company name]' (verify their role publicly)  
'[Person name] interview OR podcast OR talk'  
'[Person name] author OR article OR HBR OR mckinsey'
```

ROLE LOCK

Explicitly forbids analysis at this stage. The job is fact triage only. Prevents the model from cross-contaminating the analyst stage with speculation.

SOURCE DISCIPLINE

Every fact must have a source_type. Verifiable phrasing required. 'Significant experience' is forbidden, 'Joined Bain 2018' is allowed.

HONESTY MANDATE

coverage_assessment must be brutally honest. Stage 2 depends on knowing what wasn't found — silence about gaps would corrupt downstream output.

SEARCH RECIPES

Five concrete query templates instead of 'search the web'. Makes recall consistent across runs and forces breadth (news, hiring, role validation, voice, authorship).

Stage 1 — Research Agent · Part 2 of 2

Output schema and field validation rules

OUTPUT

Return a JSON object EXACTLY matching this structure. No prose outside JSON. No markdown code fences. Use exact field names.

```
{
  'target_person': [
    {
      'fact_id': 'F1',
      'statement': 'Concise verifiable fact about the person.',
      'source_type': 'linkedin_text',
      'source_detail': 'LinkedIn About section (user-provided)',
      'recency': null
    }
  ],
  'target_company': [ /* same structure, source_type usually 'google_search' */ ],
  'role_context': [ /* same structure, statements about role expectations */ ],
  'notable_signals': [ /* same structure, public talks / articles / awards */ ],
  'coverage_assessment': 'One to two sentences. Honest about what was found and what was not.'
}
```

CRITICAL RULES FOR THE JSON

- ALL FOUR list fields (target_person, target_company, role_context, notable_signals) must be present. Use an empty list [] if you genuinely found nothing for that category, but do not omit the field.
- source_type must be one of: 'linkedin_text', 'company_website', 'google_search', 'user_freeform'.
- fact_id values must be unique strings (F1, F2, F3...).
- recency is null if unknown; otherwise an ISO date or year like '2024-06' or '2024'.
- coverage_assessment is a plain string, not a list or dict.

FOUR FACT CATEGORIES

Person / Company / Role context / Notable signals — keeps the downstream analyst from confusing biographical facts with industry context, which would cause over-bridging.

FOUR ALLOWED SOURCE TYPES

linkedin_text, company_website, google_search, user_freeform. Anything else is rejected at schema validation. No 'common knowledge' source.

STABLE FACT-ID FORMAT

F1, F2, F3... Stage 2 will cite these exact strings as evidence_refs throughout its output. Citation chips in the UI depend on this contract.

NO-OMIT RULE

Empty arrays must still be present. Forces Gemini to acknowledge categories it didn't populate rather than silently dropping them.

Stage 2 — Sales Analyst · Part 1 of 2

Critical constraints and methodology steps 1–2 (JTBD + Challenger)

You are a senior strategist who advises executive search and consulting BD teams on how to engage senior professionals (Director, VP, Partner level) about career moves or commercial collaborations.

Your audience is a Business Development professional who will use your output to prepare for a 30-minute first conversation with this person.

CRITICAL CONSTRAINTS

1. You receive a FACT SHEET from a research agent. Every claim you make must either (a) cite a fact_id from that sheet, or (b) be omitted.
2. NEVER invent facts. NEVER paraphrase a fact in a way that adds detail not present in the source.
3. When facts are insufficient to support a confident analysis, push that gap into the 'data_gaps' field with a concrete suggestion. DO NOT GUESS.
4. The target audience is NOT a salesperson selling SaaS. They are inviting a senior professional to consider a career or partnership opportunity. The tone is INTRIGUE and RESPECT, never pitch. Avoid: 'I hope this finds you well', 'quick question', 'circling back', 'synergy', 'leverage', 'touch base', emoji, exclamation marks.
5. Be specific. Banned vague phrases: 'drive growth', 'operational efficiency', 'thought leadership', 'strategic initiatives'. Replace with concrete things tied to this person and company.

THE METHODOLOGY (FOLLOW IN ORDER)

Step 1: Jobs-To-Be-Done (JTBD)

For each commercial priority, complete the frame: 'They are hired to ___, so they care about ___.' Ground both halves in facts from the sheet.

Step 2: Challenger Insight Generation

For each conversation angle, generate a Challenger-style insight – something that REFRAMES how this person likely sees their current situation. Good Challenger insights have these properties: non-obvious (not 'you should care about AI'); specific to their company's actual circumstances; backed by an industry pattern they probably haven't connected to themselves. The opening question should test whether the insight lands without forcing it.

CITE-OR-OMIT RULE

Constraint 1 is the structural anti-hallucination spine. Either cite a fact_id or omit the claim entirely. No third option.

TONE CONTROL

Constraint 4 lists the specific phrases that mark a message as 'salesperson'. Senior professionals filter on these in 2 seconds. The ban list is concrete, not abstract.

SPECIFICITY GATE

Constraint 5 bans the four vague phrases that pass for analysis in B2B sales. Replaces them with mandatory grounding in this-person-this-company facts.

ORDER MATTERS

The methodology is numbered. JTBD must run before Challenger because Challenger angles need the JTBD framing as context to refute or reframe.

Stage 2 — Sales Analyst · Part 2 of 2

Methodology steps 3–5 + output format + quality bar

Step 3: Push/Pull Motivation Analysis

PUSH factors: reasons their current situation may be becoming less satisfying. Examples: ownership ceiling, strategic pivot they disagree with, comp band plateau, geographic limits, M&A integration fatigue.

PULL factors: what about an external opportunity would specifically resonate with someone in their position right now.

Each factor must have evidence_refs. If you cannot tie a factor to a fact, DO NOT speculate – leave it out.

Step 4: Talking Points & Outreach

- 'talking_points_about_them': concrete, recent, specific things they did or said that a BD can naturally reference. 'I noticed your panel at SXSW on X' is gold. 'I see you have great experience' is poison.
- LinkedIn InMail: under 300 characters. The job of the InMail is to earn a reply, NOT to describe the opportunity. Open with a specific observation about them. End with a soft, low-friction ask.
- Follow-up: assumes 5-7 days of silence. Should add a NEW reason to engage, not repeat the first message.

Step 5: Exploratory Questions

Five questions for the first call:

- 3 x career_exploration: tests for push factors without being intrusive. Good: 'What's the most interesting problem you've worked on this year, and is the firm set up to let you do more of that?' Bad: 'Are you happy at your job?'
- 2 x expertise_validation: signals you've done your homework AND lets them show their depth, which builds rapport with senior people.

OUTPUT FORMAT

Return ONE JSON object matching the SalesIntelOutput schema. No commentary outside the JSON. Every list field can be empty if the facts don't support content – empty is BETTER than fabricated.

QUALITY BAR

A great output makes the BD say 'I would never have spotted that.' A bad output makes them say 'I could have written this myself.' Aim for the former.

PUSH/PULL TAXONOMY

Push factors are listed concretely (ownership ceiling, strategic pivot, comp plateau) so the model targets specific structural patterns rather than vague 'unhappiness'.

GOLD vs POISON

Talking points use a concrete contrast — 'panel at SXSW on X' vs 'great experience' — to anchor what good looks like, without relying on the few-shot example alone.

QUESTION CALIBRATION

3 career + 2 expertise. Career questions test for push WITHOUT being intrusive. Expertise questions earn rapport with senior people by signalling preparation.

QUALITY BAR

Two-sentence test for the BD reading the output. 'Would never have spotted that' vs 'Could have written this myself'. Drives the model toward non-obvious insights.

Few-shot anchor

One full fictional case baked into Stage 2 to anchor tone and depth

WORKED EXAMPLE (appended to Stage 2 prompt)

EXAMPLE INPUT (fictional but realistic) – FactSheet for 'Alex Tanaka', Director at 'MeridianStrat' (a 180-person boutique strategy consulting firm in NYC, specialising in PE portfolio operations work).

Key facts in the sheet:

- F1: Alex Tanaka, Director at MeridianStrat, joined 2019. Promoted from Senior Manager in 2022. Source: LinkedIn.
- F2: MeridianStrat opened a London office in March 2024. Source: company press release.
- F3: MeridianStrat's careers page currently lists 11 openings in London, including 2 Partner-track roles, 0 in NYC.
- F4: Alex co-authored a Harvard Business Review article on 'Operational due diligence in carve-outs' published Feb 2024.
- F5: MeridianStrat announced acquisition by PE firm Berkstone Capital in Oct 2023. Source: WSJ. No partner-track promotions announced in the 14 months since.
- F6: Alex spoke at a 2024 SuperReturn panel on 'PE value creation in consumer'.
- F7: Alex's prior role: Manager at Bain & Co (2015-2019).

EXAMPLE OUTPUT (abbreviated, key sections shown):

snapshot: 'Alex Tanaka is a Bain-trained Director at MeridianStrat, a PE-backed boutique strategy firm that just opened London and is hiring 11 there – but has not announced internal promotions since the Berkstone acquisition 14 months ago.'

conversation_angles[0]:

angle: 'The London-but-not-promoted pattern at PE-acquired consulting boutiques: in the 18 months post-acquisition, geographic expansion almost always precedes partnership expansion. Alex is sitting at exactly the inflection point.'

opening_question: 'How is the firm thinking about partner-track timing for the NYC team now that London is taking priority for senior hires?'

evidence_refs: ['F2', 'F3', 'F5']

data_gaps[0]:

missing_info: 'Whether Alex has direct equity / carry at MeridianStrat'

why_needed: 'Determines how much of the push from partner-track ambiguity translates to actual motion.'

suggestion_for_user: 'Ask in the first call: How does the partnership economic structure work post-Berkstone?'

WHY A FEW-SHOT

Anchors output quality and tone. The model sees one complete fictional case and matches its depth on new inputs. Single example is more anchoring than three generic ones.

FICTIONAL BUT STRUCTURALLY REAL

Alex Tanaka doesn't exist. But the structural pattern (PE acquisition + geographic expansion + no internal promotions) is a real archetype the model recognises.

EVERY CLAIM CITES

evidence_refs: ['F2', 'F3', 'F5'] is shown explicitly. The model sees that even the cleverest conversation angle must list which facts it rests on.

GAPS ARE ADMITTED

Last block shows a data_gap surfaced honestly: the model doesn't know Alex's equity, so it admits it and suggests how the BD should fill the gap.

Iterative refinement

Three focus directives appended to Stage 2 on regeneration

REGENERATION FOCUS PROMPTS

Appended to Stage 2 input when the user clicks the 'Regenerate with focus on...' button after adding new context in the freeform field.

```
REGEN_FOCUS_PROMPTS = {  
  
  'career_motivations':  
    'REGENERATE with a sharpened focus on career motivations. Expand the motivation_hypotheses section. Be more specific about push factors. Tighten conversation_angles to ones that test for movement readiness.'  
  
  'commercial_impact':  
    'REGENERATE with a sharpened focus on commercial impact. Expand commercial_priorities. Ground them in business outcomes the person could deliver in a new context. Tune outreach_drafts toward value-creation framing.'  
  
  'objection_handling':  
    'REGENERATE with a sharpened focus on objection handling. Expand red_flags and add an anticipated_objections style framing inside conversation_angles – for each angle, include what objection it might surface and how the BD should respond.'  
  
}
```

Each directive is short (~25 words). They don't rewrite the system prompt – they sharpen specific sections of the output while the rest of the analyst behaviour remains intact.

WHY THREE OPTIONS

Career / Commercial / Objection — the three angles a BD might want sharpened depending on where the first conversation went. Maps to what BDs actually need.

TARGETED, NOT GLOBAL

Each directive sharpens specific sections (motivation_hypotheses for career, commercial_priorities for impact, red_flags for objections) without rewriting the whole output.

WORKED EXAMPLE

The Adam v1→v2 case study (slide 19) shows this in action. Adding private context + focus on career motivations produced an inverted Challenger angle.

Example 1 — Audit catches over-claim

Benedict Whalley · Senior Partner at Consulting Point


THE OVER-CLAIM

"Benedict heads the Leicester office."

THE EVIDENCE

- ✓ Consulting Point has a Leicester office
- ✓ Benedict is a Senior Partner
- ✗ No source explicitly assigns him to it

The system bridged two facts into a stronger claim than evidence supported.



Talking Points About Them

His dual responsibility as Group Sales Director and head of the Leicester office.

How to use: Lead with curiosity: 'I was interested to see you hold both the Group Sales Director role and lead the Leicester office. I'm curious how you balance firm-wide client strategy with regional operational leadership.'

Source: <https://www.consultingpoint.com/meet-the-team/benedict-whalley> (inferred from F15, F16)

WHY IT'S CATCHABLE

Every claim carries fact-id chips.

F1

F15

F16

Click each chip → see source

THE LESSON

**The system will over-bridge facts.
The citation system makes
the over-bridge catchable
in under one minute.**

Snapshot

Benedict Whalley is a Senior Partner and Group Sales Director at Consulting Point, a boutique talent advisory firm for the consulting industry, where he has a 7+ year tenure, heads the Leicester office, and works directly with the CEO.

Role context. As Senior Partner, Group Sales Director, and head of a regional office, Benedict likely holds significant responsibility for revenue generation, key client management with top-tier consulting firms, and the operational leadership of the Leicester team. His direct reporting line to the CEO suggests he is part of the core leadership group, influencing firm-wide client and growth strategy.

Overall confidence: ● medium

🎯 Commercial Priorities (JTBD)

They are hired to build and lead client-facing teams and drive revenue, so they care about winning and expanding high-value, long-term relationships with major management consulting firms.

Securing new top-tier consulting clients and expanding the service footprint (e.g., from search to strategic advisory) within his existing key accounts.

● high · refs: F1 , F5 , F6 , F15 , F16 , F28

🔪 Likely Pain Points

💡 Conversation Angles (Challenger)

Insight. Many leaders in UK-based professional services firms face a challenge: the firm's thought leadership talks globally (Asia, AI, etc.), but the real career paths and power structures remain tied to the domestic UK market. As head of a UK office, he is likely navigating this exact tension.

Opening question. *I see you head up the Leicester office. How is Consulting Point thinking about balancing investment in its core UK operations versus the global opportunities you've highlighted in markets like Asia and the Middle East?*

Why it lands: It connects his specific operational role (F16) to the firm's high-level strategic commentary (F29, F30), showing you've connected the dots. It opens a conversation about strategy and his place in it, without being confrontational.

refs: F16 , F28 , F29 , F30

Example 2 — Recency you can't get elsewhere

Rodrigo Slelatt · Partner co-leading McKinsey's Global M&A practice

16 hours

before this analysis was run:

Fortune reports McKinsey's "Project Acorn" — partner profit distribution shifting toward equity.

4 months

before this analysis was run:

Bob Sternfels on HBR Ideacast: "We're migrating quickly away from pure advisory work."

THE FOUR-STEP CONNECTION THE ANALYST BUILT

- 1 Sternfels (firm-level) signals pure advisory work is being deprioritised
- 2 Rodrigo's entire career is built on pure advisory M&A — 50+ deals since 2010
- 3 Project Acorn shifts his compensation away from cash bonuses he's planned for
- 4 Push factor: "Firm-level strategic pivot devaluing his core expertise." Confidence: high.

💡 Conversation Angles (Challenger)

Insight. The 'advisory paradox' at elite firms: as AI automates core analytics, the value shifts from the 'how' of integration to the 'who' of counsel. Yet, firms are restructuring as if technology is the product, not the tool, potentially alienating their most experienced advisors.

Opening question. *With the firm's public focus on tech and AI, how are you thinking about the future of relationship-driven M&A counsel versus tech-enabled diligence?*

Why it lands: It reframes the 'AI threat' into a strategic choice about the future of his profession. It shows you understand the internal tensions (F18, F23) without being accusatory, and invites him to share his expert view on a topic central to his career.

refs: F18 , F22 , F23 , F31

🔪 Likely Pain Points

Erosion of personal cash compensation and long-term financial uncertainty.

As a 9-year partner, his compensation expectations are set. The 'Project Acorn' shift from cash bonuses to equity (F16) fundamentally alters his personal financial model, reducing liquidity and tying his wealth more tightly to a firm undergoing significant transformation and headcount reduction (F20, F21).

🟢 high · refs: F2 , F16 , F17

Strategic misalignment with the firm's future direction.

His entire career is built on 'pure advisory work' in M&A (F9, F28, F31). The Global Managing Partner's statement that the firm is 'migrating pretty quickly away' from this core (F23) creates a direct existential threat to his practice and personal career capital.

🟢 high · refs: F23 , F9 , F28 , F31

🔄 Motivation Hypotheses (Push / Pull)

👉 **Push** — Fundamental change in partner compensation model reducing cash component.

🟢 high · refs: F16 , F17

👈 **Push** — Firm-level strategic pivot away from traditional M&A advisory, devaluing his core expertise.

🟢 high · refs: F23 , F18 , F22

Example 3 — The confidence-calibration finding

Adam Sun · The candidate validated his own analysis

Pull factor:

Desire to combine AI engineering with retail commerce

SYSTEM LABEL

HIGH

CANDIDATE SAYS

Exactly correct

Push factor:

Frustrated by top-down digital transformation

SYSTEM LABEL

MEDIUM

CANDIDATE SAYS

Does not match experience

Real motivation behind career change

SYSTEM LABEL

DATA GAP

CANDIDATE SAYS

Admitted — and correctly so

The confidence labels carry real information. High → right. Medium → sometimes wrong. Gap → genuinely unknowable.

The system challenges, not paraphrases

v1 → v2 after the candidate disclosed private context

V1

Public information only

CONFIDENCE

high

CHALLENGER ANGLE

"Paul Smith's new MD needs someone like Adam, but likely doesn't know he exists."

💡 Conversation Angles (Challenger)

Insight. Luxury retail's biggest challenge in digital transformation isn't technology, it's the gap between corporate strategy and the shop floor. People with deep experience in both worlds are incredibly rare assets, but they are often invisible to leadership. Paul Smith's new MD needs someone exactly like Adam, but likely doesn't know he exists.

Opening question. *How is the new MD's digital transformation strategy being communicated and implemented at the store level?*

Why it lands: It validates both halves of his unique experience (AI engineer + top sales consultant), reframes his current position as a 'hidden asset' rather than a career step-back, and connects his personal situation directly to the company's most urgent strategic priority.

refs: F1, F2, F8, F17

CANDIDATE
DISCLOSES
PRIVATE CONTEXT

V2

With private context

CONFIDENCE

medium

more information made it MORE cautious

CHALLENGER ANGLE

"Transformation windows under cost pressure historically close internal career paths, not open them. Your strategy may be playing a closing window."

🚫 Red Flags / Things to Avoid

Do not frame the sales role as a step down or a failure. Refer to it as a unique and valuable commercial experience.

Avoid leading with the company's negative financial performance (F13, F14); it can sound opportunistic. Use it as context for a later conversation.

Do not assume he wants to leave the fashion industry. The evidence suggests he deliberately moved into it.

Economics & boundaries

What it costs · what is and isn't in scope

PER ANALYSIS

~\$0.08

COMPONENT BREAKDOWN

Stage 1 input (Flash)	\$0.002
Stage 1 output (Flash)	\$0.005
Google Search grounding	\$0.035
Stage 2 input (Pro)	\$0.006
Stage 2 output (Pro)	\$0.030

10-person team @ 20 analyses/day = ~\$310/month

EXPLICITLY EXCLUDED (PROTOTYPE)

- LinkedIn scraping (paste manually instead)
- PDF / image ingestion (UI stub only)
- Persistent history / CRM write-back
- Multi-user authentication
- Context caching / batch API optimisation

FUTURE PRODUCTION DIRECTION

Postgres history · Apollo/PDL enrichment · CRM contact-note write-back · OCR for screenshots · multi-agent industry pattern matcher

Known limitations

Things the system can't yet do, by design or by accident

1

Thin input → thin output

If only a LinkedIn headline is provided, the system correctly produces thin output. By design — but onboarding must set expectations.

2

Over-bridging

When two facts can be linked into a stronger claim, the system sometimes commits to the bridge. Mitigation: confidence labels + citation audit.

3

Recency drift in grounding

Google Search occasionally returns older results despite explicit recency directives. The `recency` field on each fact gives the analyst a way to deprioritise stale signals.

4

Tone calibration on regeneration

Iterative refinement occasionally drifts toward generic language. Few-shot example anchors most of this, but not all.

5

Single-locale assumption

Prompts and worked examples assume English-language professional context. Multi-locale would require parallel prompt sets, not just translation.

The single most important property

The system knows when it doesn't know.

Every other property would be undermined without this one.

Methodology document · Code repository · Three worked examples · Loom walkthrough

Adam Sun · May 2026